

# COSC265 — Relational Database Systems

Neville Churcher

Department of Computer Science & Software Engineering  
University of Canterbury

2021



## Index Structures

- ★ Physical order of records on disc affects efficiency of individual queries requiring file traversal.
- ★ `select * from node where label = 'root'`
- ★ Physical records can only be in one order at a time!
- ★ Need an efficient way to access records in different orders without expensive operations like physical sorting
- ★ Index maps *values* of indexing field/attribute to corresponding logical *records*
- ★ Index entries ordered by indexing field value
- ★ Typically implemented as indexing field values with list of pointers to disc blocks containing logical records
- ★ May have multiple indices on single file

# What to Index?

Make queries go faster ...

**Primary key:** to give natural sorted order

**Foreign key:** to assist in join operations

```
select name from node, edge
  where edge.from = node.id
```

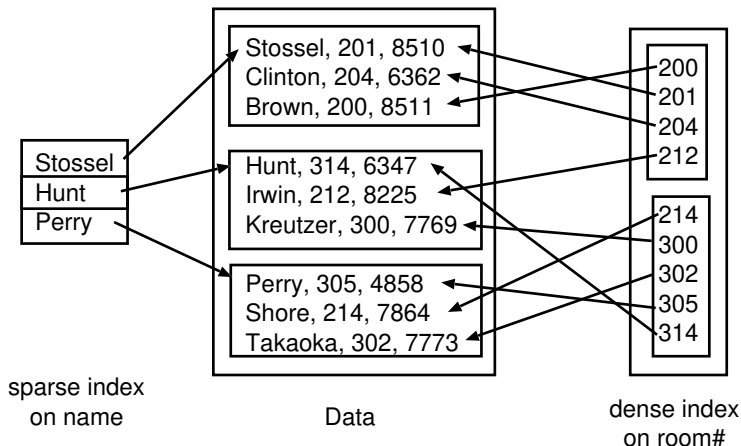
**Common access paths:**

- ★ determined at design time or
  - ★ observed in use (tuning)
  - ★ may involve attributes that aren't in PKs or FKs
- ```
select * from student where gpa < 1
```

## Classifying Indices

**Dense:** at least one index entry for each search key value


**Sparse:** index entries for only some search key values (typically anchor records for primary index)



## Classifying Indices

**Primary:** includes primary key field(s). Physical order determined by unique, non-null, *ordering key field*

**Secondary:** may contain duplicates. Implementation may use block or record pointers



## Primary Index

★ (Ideally) Maintains entity integrity constraint

★ Spot the difference...

```
create table bar (k int, primary key(k))
```

```
create table bar (k int not null)
```

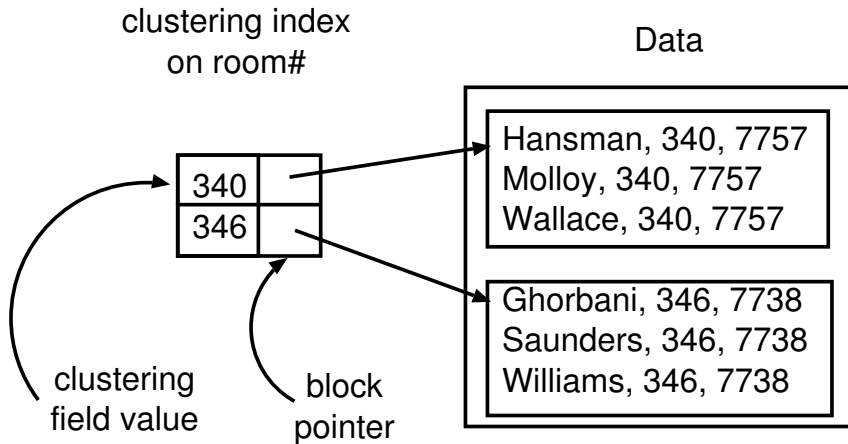
```
create table bar (k int)
```

```
create unique index ki on bar(k)
```

# Clustering

- ☆ Primary index specified on ordering key field.
- ☆ If ordering field is not a key (i.e. can have duplicates) then clustering index can be used
- ☆ clustering — *logically related* items placed *physically close* on disc
- ☆ Since a file can have at most one ordering field, can have at most one primary index or one clustering index *but not both*
- ☆ Clustering index is sparse—one entry per value rather than per record

## Clustering Example





## For really big data sets

- ☆ Index entries ordered by indexing field value
- ☆ Index files can become very large (though usually smaller than data file)
- ☆ Index traversed/searched often (otherwise pointless!)
- ☆ Single-level uses index binary search ( $\log_2 N$ )
- ☆ Multilevel indexing involves replacing simple (single-level) index with more complex data structure
- ☆ Time-space trade-offs
- ☆ Evolution factors—index entries may change  $\nleftrightarrow$  data record updates

## Index Implementation

Single-level: uses index binary search ( $\mathcal{O}(\log_2 N)$ )

Tree structures: speed up searching

- ☆ Internal nodes contain many pointers to other nodes
- ☆ Leaf nodes contain key values and pointers to physical data
- ☆ Balanced trees preferable when index changes often. Node split if contains less than  $M/2$  keys
- ☆ Common examples include B-trees, B+-trees (no internal nodes have data pointers) and variants
- ☆ Searching order  $M$  B-tree  $\mathcal{O}(\log_{M/2} N)$  — effectively constant if  $M$  large enough

## Multiple Attribute Indexing

- ★ Sometimes want to index on combination of attributes (e.g. Surname + GivenName)
- ★ Could use single index and select from results
- ★ Often major and minor component
- ★ May be better to create *surrogate* attribute for combination if need more than two attributes in index

## Indexing in Oracle

- ★ Btree is default
- ★ Some other types available (see docs)
- ★ If the `create table` DDL statement has a **primary key** clause then Oracle will automatically create an index for the PK attributes
- ★ However, Oracle — unlike many other DBMS — will *not* automatically create an index if there is a **foreign key** clause
- ★ `create index by_name on customer(surname)`
- ★ `create index by_location on supplier(country, city)`

## Index Maintenance

- ★ Costs associated with updates to index as well as data
- ★ Insertion & deletion require moving other records as well as updating index entries
- ★ Insertion/Deletion for primary index complicated by changes to anchor records
- ★ DBMS implementations utilise many variations
- ★ For bulk loading etc. it may be more efficient to drop indexes, load data then rebuild indexes

## Design Issues

- ★ Tradeoff time/space/performance
- ★ Tuning data (statistics, histogram, ...)
- ★ Query optimisation
- ★ Predictability of queries (PK/FK, ...)
- ★ DBMS may silently create its own indexes for us
- ★ DBMS may have special purpose index types (e.g. quadrees for spatial indexing), inverted indexes for full text searching, ...